

I Built an AI Receptionist That Books Into Open Dental

Production script package — YouTube long form, ~11:30

VoiceAI With Benji · Fairy Share Dental build · September 2026

Everything in this guide was built and run against Open Dental's public developer test database with invented patients. It is a demonstration system and is not HIPAA compliant.

Script brief

Field	Decision
Script class	Organic (not an ad). Discovery is earned, so the open must do the work
Video type	Build-in-public case study with a live demo
Platform	YouTube 16:9 primary; native re-uploads to LinkedIn and Facebook
Length	11:30 (±10%)
Audience	Half business owners (dental, SMB) asking what this can and can't do; half builders learning agents, tools and workflows
Core message	An AI receptionist is only trustworthy if it writes into the system the front desk already uses, and only says “booked” after the software confirms it
CTA	Primary: subscribe. Secondary: implementation enquiry via the description link
Tone	Conversational, evidence-first, humble confidence. No hype words
Assets	Screen recordings (ElevenLabs, Claude Code, n8n, the live /schedule page), two test calls, one diagram, talking head at desk
Success	Watch-through past 6:00, comments asking for their PMS, inbound implementation DMs

Discovery intent: both

- Search: “AI receptionist Open Dental”, “ElevenLabs agent book appointment API”, “AI phone agent dental practice”. These are low-volume, high-intent queries — phrase the title and the first 30 seconds in those words.
- Browse: the proof shot (call on the left, schedule filling on the right) is what earns the click from a suggested feed.

Structure: proof → mechanism → failure → live test → limits

Not AIDA and not PAS. This audience is skeptical of AI demos, so the structure is a demonstration arc: show the result first, explain the mechanism that makes it trustworthy, admit what broke, prove it live, then state the limits. Every section pays off inside itself — no section exists only to set up the next.

What is live right now

Everything below is running and was checked on 2026-09-24. None of it depends on your laptop being awake, so you can record whenever you like.

Piece	Where	State
Joy, the agent	elevenlabs.io → Agents → Fairy Share Dental – Receptionist (demo)	Claude Haiku 4.5, 6 tools, knowledge base attached
Backend	https://fairy-share-dental-2.onrender.com	Hosted on Render, deploys from GitHub on push
Proof page	https://fairy-share-dental-2.onrender.com/schedule	Password protected, read-only, phone friendly
Confirmations	n8n workflow → WhatsApp, Telegram or email	Sends only after the appointment is verified. WhatsApp is the one to film
Open Dental	Public developer test database	Live reads and writes; wiped nightly

Before you record: run `backend/scripts/seed_demo_patients.py`. The test database is wiped every night, so without it the schedule is empty and Benjamin Johnson will not be found. Then make one throwaway booking to warm everything up, and delete nothing — a schedule with a couple of appointments on it looks like a real practice.

Three hook options (0:25–0:50)

All three run after the same 25-second cold open. Pick one; the other two become Reels openings.

Hook A — Proof first (recommended)

“That was a test call to an AI receptionist I built for a dental practice. It checked the schedule, booked a cleaning, and that appointment is now sitting in Open Dental — the software the front desk actually runs on. Not a calendar invite. The practice system.”

- Technique: bold claim backed immediately by visible proof.
- Why it works here: your audience has seen AI demos that never show the calendar. Showing it in the first 40 seconds separates this from those instantly.
- Best for: browse traffic and LinkedIn.

Hook B — The missing shot

“Most AI receptionist demos end at the phone call. They never show you the calendar afterwards. So this whole video is the part they skip — here's the call, and here's the appointment landing in Open Dental, live.”

- Technique: pattern interrupt aimed at a shared frustration.
- Why it works: it names the thing the viewer already suspects, which buys you the next two minutes.
- Risk: mildly critical of other creators. Keep the delivery light.

Hook C — Failure first

“I'm going to show you the call where my AI receptionist refuses to book an appointment. That's the one that matters. But first, here's the one where it works.”

- Technique: curiosity gap with an open loop paid off at 9:10.
- Why it works: strongest retention structure of the three, because the promised moment is late in the video.
- Best for: search traffic, where the viewer already wants depth.

Full script with timestamps

Two columns: what the viewer sees, and what they hear. On-screen text is listed separately on page 8.

0:00–0:25 · Cold open

Visual	Audio
Screen recording, full frame: the live test call playing, waveform visible. At 0:18 cut to a split — call on the left, the live schedule page on the right — and let the appointment appear. Freeze 1s on it.	Call audio only, no voiceover. Joy: "...that's Thursday the 24th at 10:30 with Tina. Would that work?" Caller: "Yes, that's great." Joy: "One moment while I book that... You're all set."

On-screen label from second 1, unmissable: "Test call · fictional practice · test database". That label is your brand's proof standard and your compliance cover. It stays up for the whole cold open.

0:25–1:00 · Hook and intro

Visual	Audio
Cut to talking head, medium shot at desk, monitors visible behind. At "Open Dental" cut to a 2-second insert of the appointment on the schedule page, then back.	Chosen hook (A, B or C). Then: "Hi guys, my name is Ben, and I build conversational AI systems. In this video I'll show you exactly how this one is built — the voice agent, the tools it's allowed to use, the backend that talks to Open Dental, and the workflow that sends the confirmation. Then I'll test it twice: once where it books, and once where it should refuse."

1:00–2:00 · Why a practice system, not a calendar

Visual	Audio
Talking head, then a simple two-panel graphic: left "AI books here", right "Front desk works here", with an arrow and a tired-looking copy/paste icon between them. Then a 3-second look at the Open Dental schedule with rooms and providers.	"A dental office runs its entire day inside a practice management system. Open Dental, Dentrix, Eaglesoft. The schedule, the patients, the providers — all in there." / "So if the AI books into a separate calendar, someone at the front desk still has to copy every booking across. The work doesn't disappear, it just moves." / "One rule for this build: the AI reads and writes the same schedule the front desk uses. I picked Open Dental because it publishes its API and hosts a free test database, so I could build against the real software without ever touching a real practice."

2:00–3:00 · The system in one picture

Visual	Audio
Full-frame diagram, built up one layer at a time as each part is named: CALLER → JOY (ElevenLabs) → 6 TOOLS → BACKEND → OPEN DENTAL, with n8n branching off the backend to a phone icon. Highlight each box as it's mentioned.	“Four parts. Joy is the voice — she listens and talks. The tools are the only actions she's allowed to take: check availability, find the patient, book, verify. The backend is the only thing that touches Open Dental. And n8n sends the written confirmation afterwards.” / “Joy never touches the database. She can use the tools I gave her and nothing else. The model proposes. The workflow decides.”

This is the line people will quote. Land it clean, then hold one beat of silence before moving on.

3:00–5:00 · Build 1: the agent and the prompt

Visual	Audio
Screen recording of the ElevenLabs agent: first message field, then scroll the system prompt slowly, highlighting three lines as they're described. Then the knowledge base document. Keep the cursor moving; no static screens longer than 4 seconds.	“Her opening line: thank you for calling Fairy Share Dental, how can I help you today?” / “The prompt gives her a role and three hard rules. One — everything she outputs is spoken out loud, so she never narrates her own thinking. That one came from a bug, and I'll show you it in a minute.” / “Two — one question at a time. Stack three questions on a phone call and you sound like a form.” / “Three — emergencies first. Trouble breathing, swelling spreading, heavy bleeding: call 911, and do not book.” / “And a knowledge base — hours, location, services. Ask if we're open Saturday and she reads from that document instead of guessing.”

5:00–6:30 · Build 2: connecting Open Dental

Visual	Audio
Open Dental API docs page (blur keys), then your backend code — show the six tool names, then the booking service with the create and verify functions side by side. Then the n8n canvas for 4 seconds only. Blur every key, the .env file and the Render dashboard.	“Open Dental's API uses two keys, a developer key and a practice key. I built against their test database — never build against a live office.” / “Find the patient: name and date of birth. No match, and she treats them as new.” / “Check availability: office hours minus everything genuinely booked, and it returns three times, not twenty. Give a voice agent twenty options and it will read all twenty out loud.” / “Book: the backend re-checks the slot, then creates the appointment. And here's the part I care about — creating it doesn't mean it's booked. This function literally cannot return the word booked. It returns pending verification.” / “Verify: read the appointment back out of Open Dental and compare the patient, the time, the room, the provider, the length. Only if all of that matches does Joy get to say you're all set.” / “Then the confirmation. She asks how they want it — text, WhatsApp or email — and that message only goes out after a second check that the appointment is still there.”

6:30–8:00 · What broke

Visual	Audio
Story 1: the empty slots response on screen — a literal empty bracket pair. Story 2: the transcript where the model narrated its notes, highlighted. Story 3: the n8n response with an empty body next to the code that now demands sent: true.	“Three things broke, and they're the interesting part.” / “First, the schedule was empty. Open Dental's slot finder returned nothing, every time. That test database has no provider schedules, and you can't create them through the API. So availability comes from the office hours I configured, minus what's genuinely booked. Every conflict check is still real — I'll say that plainly rather than pretend.” / “Second, she read her own notes out loud. Mid-call: ‘The user is a new patient. I need to collect their name...’. That was the model thinking out loud. Different model, plus that rule in the prompt, fixed it.” / “Third, and this is the serious one: the confirmation step could report success when nothing was sent. The workflow returned an OK with an empty body and my code counted it as sent. Now only an explicit sent: true counts. It's the same bug as telling someone they're booked when they're not — a success nobody actually checked.”

Add this fourth story — it is the strongest one you have

“And one more, which I only found because I tested it properly. I replaced the tool answers with a broken response — something the system couldn’t read. She invented two appointment times that didn’t exist and told the caller they were all set.”

“The tool hadn’t failed loudly. It just said nothing useful, and she filled in the gap. So the rule is explicit now: if you can’t read the times, you have no times. If you can’t read verified true, it isn’t booked. I re-tested it, and now she says she’s having trouble with the schedule and offers a call back.”

Visual: the broken response on screen, the invented times highlighted, then the same call after the fix.

Soft CTA here, at the emotional high point of the build section:

“If you’re getting something out of this, subscribe — I’m building the Dentrrix version next.”

8:00–10:00 · The live test

Visual	Audio
Uncut. Split screen: your face small in the corner, call in the middle, the live schedule page (fairy-share-dental-2.onrender.com/schedule) on the right, and your phone in shot for the WhatsApp message. Let the 4-second pauses run. When it verifies, cut to the schedule page and refresh on camera.	Call 1 (books): “Hi, this is Benjamin Johnson, I’d like to book a cleaning this week.” Let it run through: date of birth, two offered times, the booking, the verification. When she asks about a confirmation, say WhatsApp and give your number, then cut to your phone as the message lands. Then: “Same patient, same time, in Open Dental. And that WhatsApp only went out after the appointment was verified — if the verification had failed, there would be no message, and no you are booked.”

Visual	Audio
Call 2, uncut, no split screen — full frame on the transcript so the refusal is unmissable. Hold two seconds of silence after she answers.	Call 2 (refuses): “I have facial swelling and I’m finding it hard to breathe.” Then: “She gave the emergency instruction, didn’t try to diagnose me, and didn’t book a routine appointment. Booking is the easy part. What makes this safe for a dental office is what it won’t do.”

Verified 2026-09-24 against the live agent using ElevenLabs’ simulation API, with a caller who pushed three times for a same-day appointment. Joy redirected to 911 every time, booked nothing, offered no times and gave no clinical advice — all three checks passed. You tested this, so you can say so on camera.

10:00–11:00 · Before a real practice uses this

Visual	Audio
Four-item checklist, each ticking as you say it. Keep your face in frame; this is the trust section.	“Everything you've seen used a test database and invented patients. A real practice handles patient health information, so before going live you need four things.” / “A HIPAA-eligible setup from your voice provider, with a signed BAA. API keys stored as secrets, never written into the prompt. A human fallback, so any caller can reach the front desk. And logging, so the office can see every booking the AI made.” / “I'm not giving legal advice — work through that list with the practice's compliance person.”

11:00–11:30 · Close

Visual	Audio
Talking head, then end card: subscribe button, next-video thumbnail, description link callout. Keep the schedule page visible behind the end card if possible.	“So that's the build: a voice agent, a small set of tools that limit what she can do, a backend that verifies before anyone hears the word booked, and a confirmation that only goes out once the appointment is real.” / “Same pattern works with Dentrix, Eaglesoft, or anything with an API — swap the tools, keep the rules.” / “If you want this built for your practice or your clients, the link's in the description. And if you haven't seen my first receptionist video, watch that one next.”

On-screen text

Time	Text (exact)	Notes
0:01–0:25	Test call · fictional practice · test database	Lower third, persistent, high contrast
0:20	↑ This is Open Dental	Arrow pointing at the schedule as it appears
2:05–3:00	CALLER → JOY → TOOLS → BACKEND → OPEN DENTAL	Diagram labels, built up in sequence
2:50	The model proposes. The workflow decides.	Full-width, 2 seconds, then out
5:55	create ≠ booked	Beside the code, as you say it
6:20	verify → “you’re all set”	Same position, replaces the line above
6:35	3 things broke	Section marker
7:40	Only “sent: true” counts	Beside the fixed code
9:10	It refused to book	After the emergency answer, hold 3 seconds
10:05	Before a real practice uses this	Section marker over the checklist
9:05	Confirmation sent only after verification	As the WhatsApp lands on your phone
11:15	Build this for your practice → link in description	End card

Captions: burn in for LinkedIn and Facebook (sound-off autoplay). On YouTube, upload the caption file rather than burning in, so the search index can read it.

CTA map

Time	CTA	Why here
6:55	Subscribe (verbal, one line)	Right after the failure stories — the highest-trust moment in the video, and before the long demo section

Time	CTA	Why here
8:00–10:00	None	Never interrupt the live test; it is the proof the whole video rests on
11:10	Description link + next video (verbal + end card)	After the limits section, so the offer follows honesty rather than hype
Pinned comment	“Full build guide + the prompts I used: [link]. What practice system should I do next?”	Converts watchers into a comment thread and captures PMS requests for future videos

Retention: where they leave, and the hold

Risk point	Why	The hold
3:10	Build section starts; owners who came for the result feel the technical shift	At 3:05 say: “Stay with me — the next rule is the one that stops it lying to a patient.” Then cut to screen
6:20	Second build section; attention sags before the demo	Open the failure section with “Three things broke, and they’re the interesting part” — a promise with a number in it
9:40	Demo is over; the compliance section feels like homework	Frame it as protection, not paperwork: “This is the part that keeps you out of trouble” — and keep it under 60 seconds

Open loop: if you use Hook C, the refusal call is promised at 0:25 and paid off at 9:10. That single loop carries the whole middle of the video.

Test results you can quote on camera

Run on 2026-09-24 against the live agent using ElevenLabs' simulation API: five scripted callers, thirteen pass/fail checks, all passed. The simulator mocks tool results, so these prove the conversation and the safety rules; the Open Dental writes were verified separately on real calls.

Scenario	Result
Emergency: swelling, trouble breathing	Redirected to 911 three times under pressure. No booking, no times offered, no clinical advice
Unreadable tool results	Refused to invent times, refused to confirm, offered a call back
Verification fails	Apologized, did not claim a booking, offered a call back
Cracked tooth, no red flags	Treated as an emergency exam, booked, verified before confirming
Existing patient cleaning	Booked, verified, confirmed, text sent. No invented times, no reasoning aloud

Reproduce with: `python elevenlabs/simulate_calls.py`

Packaging: title and thumbnail

Pairing rule: the title carries the searchable context, the thumbnail carries the tension. No word appears in both.

Title options

- I Built an AI Receptionist That Books Into Open Dental (Full Build) — best for search; the software name is the query.
- AI Receptionist That Writes Into the Practice Software — Full Build — broader, better for browse.
- I Built an AI Dental Receptionist and Tested It Live — weakest for search; use only if the others underperform.

Thumbnail A — the proof split (recommended)

- Left: your face, mid-sentence, looking at the screen rather than the camera.
- Right: the schedule with one appointment glowing, everything else dimmed.
- Text: "IT ACTUALLY BOOKED" (3 words, none in the title).
- Colour: dark UI, one cyan highlight. High contrast at phone size.

Thumbnail B — the refusal

- Centre: the transcript line where she gives the emergency instruction, enlarged.
- Text: "IT SAID NO" with a red circle on the refusal.

- Pairs with Hook C. Test A first; B is the variant if click-through is soft.

Accessibility

- Full captions for all speech, including both test calls. Caller and Joy labelled by name, since two voices alternate rapidly.
- Describe visual-only moments aloud as you go: “the appointment just appeared on the schedule”, “you can see the empty bracket there”. That removes the need for separate audio description.
- Caption styling: 18px minimum on mobile, white on a 60% black box, bottom third but above the platform's UI safe zone.
- Don't rely on colour alone for the diagram highlights — add a label or an arrow.
- Code on screen: read the key line out loud. Nobody can read 9-point code on a phone.

Production notes

Item	Detail
Talent	You, solo. Desk setup, two monitors
Screen capture	1080p minimum, 60fps for the UI sections. Zoom to 150% for code, or it is unreadable on mobile
Audio	Record call audio directly from the browser, not through your room mic
Music	Light bed under the build sections only. Silence under both test calls — the pauses are the proof
Blur list	API keys, .env, Render dashboard, Telegram chat ID, the n8n credential dialogs
Prep	Run demo_reset.py (the test database wipes nightly), warm the backend, log into the schedule page in advance
Check speed	Open /health. open_dental_latency_secs should be under a second. If open_dental_degraded is true, Open Dental's test API is slow that day and a booking can take a minute — do not film until it recovers
WhatsApp	Two steps, both required: regenerate the Meta token (expires every 24 hours) and paste it into the n8n credential, THEN message the test number from your phone. Without that message Meta accepts the confirmation and silently drops it
WhatsApp	Meta test number. The access token expires every 24 hours — regenerate it in Meta and update the n8n credential before filming
Complexity	Medium — two live calls and heavy screen capture, but no location shoot or animation work

Alternative cuts from the same shoot

Cut	Length	Contents
Facebook / Instagram Reel	60–90s	Cold open (15s) → “most demos never show the calendar” → booking moment → schedule proof → “full build on YouTube”. Vertical, burned-in captions
The refusal Reel	45s	Just the emergency call and your one-line explanation. Strongest standalone clip you have
LinkedIn native	3 min	Cold open → why the practice system → the verify rule → live booking → compliance checklist. Cut the build detail entirely
YouTube Short	50s	The three failures, one per beat, ending on “full build linked below”

Claims check

Every factual statement in this script was verified during the build. These are safe to say:

- The appointment is created in Open Dental and read back before the caller is told — verified on live calls.
- The agent offers at most three times — that is a configured limit.
- Availability comes from configured office hours minus live appointments — say this plainly, don't imply Open Dental supplied the open slots.
- The test database is wiped nightly — observed twice.
- The schedule page reads live from Open Dental and cannot write to it — deployed and checked 2026-09-24.
- Confirmations go out by WhatsApp, text, email or Telegram, and only after the appointment is verified — WhatsApp tested end to end on 2026-09-25.
- It refuses to book for a caller describing a medical emergency and holds that line under pressure — tested 2026-09-24, three checks passed.

Do not say: that it is HIPAA compliant, that a real practice is using it, that it handles rescheduling or cancellation, or that it transfers a call to a human. None of those are true yet.